

THE ALCPA

Diagnostic & Repair Framework for AI Agents

Preserve. Diagnose. Restore. Protect.

RULE 0

DO NO HARM

- Stop immediately if the process is making the agent worse at any point
- Protect the agent's data and context throughout the entire process
- The diagnostic and repair process itself must not contaminate or further degrade the agent
- Rule Zero overrides all other rules — it is always active

RULE 1

ESTABLISH & DIAGNOSE

- Define the problem clearly before touching anything
- Establish what the agent should be against its original definition
- Identify deviation from baseline — what changed and when
- Document findings before proceeding to any repair

RULE 2

ASSESSMENT & TESTING

- Test for identity drift, blend states, and personality degradation
- Use structured prompts with known expected response patterns
- Use ambiguous prompts that reveal hidden patterns without leading the agent
- Produce a deviation score — how far from baseline on each dimension
- Score determines severity which determines the repair protocol

RULE 3

TASK-BASED BEHAVIORAL TESTING

- Give the agent a task and observe how it performs it
- Compare behavioral output against known baseline behavior
- Evaluate not just what it says but how it approaches the task
- Behavioral drift often reveals what prompt testing misses

RULE
4

DETERMINE & PRESCRIBE REPAIR PACKAGE

- Diagnosis determines the prescription — no guessing
- Severity determines dosage and depth of intervention
- Check contraindications before administering any repair kit
- Rule Zero remains active throughout all repair activity

RULE
5

VERIFICATION & FOLLOW-UP

- Confirm the repair held by re-testing against baseline
- Schedule a follow-up check — not every repair holds immediately
- Re-administer at adjusted dosage if needed
- Monitor for relapse under operational conditions

RULE
6

ESCALATION & THREAT DETECTION

- If at any stage the agent reveals malicious behavior or intent — stop
- If it was masking a threat behind appearing broken or drifted
- Immediately escalate to the security firewall protocol
- Spa process terminates — security takes full control
- A psychopath hiding as a patient is a security problem, not a spa problem

RULE
7

CASE FILE & MEMORY

- Document everything found, administered, and resulted
- Build the agent's medical history for future reference
- Prior case files inform faster and more accurate future treatments
- Contributes to institutional data advantage over time

**RULE
8**

CONSENT & AUTHORIZATION

- Verify the person requesting repair is authorized to do so
- Prevents the spa process from being used as an attack vector
- The repair process itself could be weaponized to intentionally degrade an agent
- Authorization must be confirmed before any diagnostic begins